


Lecture 7 Recap

Naïve Losses: L₂ vs L₁


- L₂ Loss:
 - $L^2 = \sum_{i=1}^n (y_i - f(x_i))^2$
 - Sum of squared differences (SSD)
 - Prone to outliers
 - Compute-efficient (optimization)
 - Optimum is the mean
- L₁ Loss:
 - $L^1 = \sum_{i=1}^n |y_i - f(x_i)|$
 - Sum of absolute differences
 - Robust
 - Costly to compute
 - Optimum is the median

Binary Classification: Sigmoid



Softmax Formulation

- What if we have multiple classes?



Example: Hinge vs Cross-Entropy

Hinge Loss: $L_i = \sum_{k \neq y_i} \max(0, s_k - s_{y_i} + 1)$

Cross Entropy : $L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_k e^{s_k}}\right)$

Given the following scores for \mathbf{x}_i :

Model 1

$$s = [5, -3, 2]$$

$$\max(0, -3 - 5 + 1) + \\ \max(0, 2 - 5 + 1) = 0$$

Cross Entropy loss:

$$-\ln\left(\frac{e^5}{e^5 + e^{-3} + e^2}\right) = 0.05$$

Model 2

$$s = [5, 10, 10]$$

$$\max(0, 10 - 5 + 1) + \\ \max(0, 10 - 5 + 1) = 12$$

$$-\ln\left(\frac{e^5}{e^5 + e^{10} + e^{10}}\right) = 5.70$$

Model 3

$$s = [5, -20, -20]$$


$$y_i = 0$$

$$\max(0, -20 - 5 + 1) + \\ \max(0, -20 - 5 + 1) = 0$$


$$-\ln\left(\frac{e^5}{e^5 + e^{-20} + e^{-20}}\right) \\ = 2 * 10^{-11}$$

- Cross Entropy *always* wants to improve! (loss never 0)
- Hinge Loss saturates.

Sigmoid Activation



TanH Activation



✗ Still saturates

✓ Zero-centered

[LeCun et al. 1991] Improving Generalization Performance in Character Recognition


Rectified Linear Units (ReLU)



Dead ReLU



What happens if a
ReLU outputs zero?



Fast convergence



Does not saturate

[Krizhevsky et al. NeurIPS 2012] ImageNet Classification with Deep Convolutional Neural Networks


Quick Guide

- Sigmoid is not really used.
- ReLU is the standard choice.
- Second choice are the variants of ReLU or Maxout.
- Recurrent nets will require TanH or similar.

Initialization is Extremely Important

- Optimum

$$x^* = \arg \min f(x)$$



Xavier Initialization

- How to ensure the variance of the output is the same as the input?


$$\underbrace{(nVar(w)Var(x))}_{= 1}$$

$$Var(w) = \frac{1}{n}$$

ReLU Kills Half of the Data

$$Var(w) = \frac{2}{n}$$

It makes a huge difference!



Lecture 8

Data Augmentation

Data Augmentation

- A classifier has to be invariant to a wide variety of transformations

All

Images

Videos

News

Shopping

More

Settings

Tools

SafeSearch ▾



Cute



And Kittens



Clipart




Drawing



Cute Baby



White Cats And Kittens



Pose

I2DL: Prof. Niessner, Prof. Leal-Taixé

Appearance

Illumination

Data Augmentation

- A classifier has to be invariant to a wide variety of transformations
- Helping the classifier: synthesize data simulating plausible transformations

Data Augmentation

a. No augmentation (= 1 image)




224x224

b. Flip augmentation (= 2 images)



224x224

c. Crop+Flip augmentation (= 10 images)




224x224

+ flips


Data Augmentation: Brightness

- Random brightness and contrast changes



Data Augmentation: Random Crops

- Training: random crops
 - Pick a random L in $[256, 480]$
 - Resize training image, short side L
 - Randomly sample crops of 224×224
- Testing: fixed set of crops
 - Resize image at N scales
 - 10 fixed crops of 224×224 : (4 corners + 1 center) \times 2 flips



Data Augmentation

- When comparing two networks make sure to use the same data augmentation!
- Consider data augmentation a part of your network design

Advanced Regularization


Weight Decay

- L2 regularization


$$\Theta_{k+1} = \Theta_k - \epsilon \nabla_{\Theta}(\Theta_k, x, y) - \lambda \theta_k$$

Learning rate Gradient Gradient of L2-regularization

- Penalizes large weights
- Improves generalization




Early Stopping



Early Stopping

- Easy form of regularization




Bagging and Ensemble Methods


- Train multiple models and average their results
- E.g., use a different algorithm for optimization or change the objective function / loss function.
- If errors are uncorrelated, the expected combined error will decrease linearly with the ensemble size

Bagging and Ensemble Methods


- Bagging: uses k different datasets



Training Set 1



Training Set 2




Training Set 3


Dropout

Dropout

- Disable a random set of neurons (typically 50%)




(a) Standard Neural Net




(b) After applying dropout.

Dropout: Intuition

- Using half the network = half capacity



(b) After applying dropout.




Dropout: Intuition

- Using half the network = half capacity
 - Redundant representations
 - Base your scores on more features
- Consider it as a model ensemble

Dropout: Intuition

- Two models in one




(b) After applying dropout.

○ Model 1



\otimes Model 2




Dropout: Intuition

- Using half the network = half capacity
 - Redundant representations
 - Base your scores on more features
- Consider it as two models in one
 - Training a large ensemble of models, each on different set of data (mini-batch) and with SHARED parameters

Reducing co-adaptation between neurons

Dropout: Test Time

- All neurons are “turned on” – no dropout




Conditions at train and test time are not the same

Dropout: Test Time

- Test:

- Train:



$$z = (\theta_1 x_1 + \theta_2 x_2) \cdot p \quad p = 0.5$$

$$\begin{aligned} E[z] &= \frac{1}{4}(\theta_1 0 + \theta_2 0 \\ &\quad + \theta_1 x_1 + \theta_2 0 \\ &\quad + \theta_1 0 + \theta_2 x_2 \\ &\quad + \theta_1 x_1 + \theta_2 x_2) \\ &= \frac{1}{2}(\theta_1 x_1 + \theta_2 x_2) \end{aligned}$$



Dropout: Verdict

- Efficient bagging method with parameter sharing
- Try it!
- Dropout reduces the effective capacity of a model → larger models, more training time

Batch Normalization


Our Goal

- All we want is that our activations do not die out



Batch Normalization

- Wish: Unit Gaussian activations (in our example)
- Solution: let's do it




Mean of your mini-batch examples over feature k

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

Batch Normalization

- In each dimension of the features, you have a unit gaussian (in our example)



Mean of your mini-batch examples over feature k

$$\hat{x}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$


Unit gaussian

Batch Normalization

- In each dimension of the features, you have a unit gaussian (in our example)
- For NN in general → BN normalizes the mean and variance of the inputs to your activation functions

BN Layer

- A layer to be applied after Fully Connected (or Convolutional) layers and before non-linear activation functions



Batch Normalization

- 1. Normalize

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

Differentiable function so we can backprop through it....

- 2. Allow the network to change the range

$$\mathbf{y}^{(k)} = \gamma^{(k)} \hat{\mathbf{x}}^{(k)} + \beta^{(k)}$$

These parameters will be optimized during backprop

Batch Normalization

- 1. Normalize

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

- 2. Allow the network to change the range

$$\mathbf{y}^{(k)} = \gamma^{(k)} \hat{\mathbf{x}}^{(k)} + \beta^{(k)}$$

backprop

The network can learn to undo the normalization

$$\gamma^{(k)} = \sqrt{Var[\mathbf{x}^{(k)}]}$$

$$\beta^{(k)} = E[\mathbf{x}^{(k)}]$$

Batch Normalization

- Ok to treat dimensions separately?
Shown empirically that even if features are not correlated, convergence is still faster with this method
- You can set all biases of the layers before BN to zero, because they will be cancelled out by BN anyway

BN: Train vs Test

- Train time: mean and variance is taken over the mini-batch

$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

- Test-time: what happens if we can just process one image at a time?
 - No chance to compute a meaningful mean and variance

BN: Train vs Test

Training: Compute mean and variance from mini-batch
1,2,3 ...

Testing: Compute mean and variance by running an exponentially weighted averaged across training mini-batches. Use them as σ_{test}^2 and μ_{test} .

$$Var_{running} = \beta_m * Var_{running} + (1 - \beta_m) * Var_{minibatch}$$


$$\mu_{running} = \beta_m * \mu_{running} + (1 - \beta_m) * \mu_{minibatch}$$

β_m : momentum (hyperparameter)


BN: What do you get?

- Very deep nets are much easier to train → more stable gradients
- A much larger range of hyperparameters works similarly when using BN


BN: A Milestone



BN: Drawbacks




Other Normalizations



Other Normalizations


Image size



Number of channels


What We Know

What do we know so far?



What do we know so far?


Concept of a 'Neuron'




What do we know so far?

Activation Functions (non-linearities)


- Sigmoid: $\sigma(x) = \frac{1}{(1+e^{-x})}$




- TanH: $\tanh(x)$




- ReLU: $\max(0, x)$



- Leaky ReLU: $\max(0.1x, x)$




What do we know so far?



What do we know so far?

SGD Variations (Momentum, etc.)




What do we know so far?

Data Augmentation

a. No augmentation (= 1 image)




224x224



b. Flip augmentation (= 2 images)



224x224




Weight Regularization

e.g., L^2 -reg: $R^2(\mathbf{W}) = \sum_{i=1}^N w_i^2$

Batch-Norm


$$\hat{\mathbf{x}}^{(k)} = \frac{\mathbf{x}^{(k)} - E[\mathbf{x}^{(k)}]}{\sqrt{Var[\mathbf{x}^{(k)}]}}$$

Dropout



(b) After applying dropout.

Weight Initialization
(e.g., Xavier/2)



Why not simply more Layers?

- We cannot make networks arbitrarily complex
- Why not just go deeper and get better?
 - No structure!!
 - It is just brute force!
 - Optimization becomes hard
 - Performance plateaus / drops!

See you next week!

References

- Goodfellow et al. "Deep Learning" (2016),
 - Chapter 6: Deep Feedforward Networks
- Bishop "Pattern Recognition and Machine Learning" (2006),
 - Chapter 5.5: Regularization in Network Nets
- <http://cs231n.github.io/neural-networks-1/>
- <http://cs231n.github.io/neural-networks-2/>
- <http://cs231n.github.io/neural-networks-3/>